



中国科学技术大学
University of Science and Technology of China

011174.01: Operating System 操作系统原理与设计

Chapter 4 之 real-time CPU scheduling

陈香兰(xlanchen@ustc.edu.cn)

高能效智能计算实验室, CS, USTC @ 合肥

嵌入式系统实验室, CS, USTC @ 苏州



温馨提示:



为了您和他人的工作学习，
请在课堂上**关机或静音**。


不要在课堂上接打电话。

- Some special issues for real-time CPU scheduling
 - Performance: How to minimize the latencies
 - Event-driven & event latency
 - RTOS
 - Priority Scheduling with preemption
 - Short interrupt latency & short dispatch latency
 - Schedulability
 - Process modeling
 - Admission-control
 - Rate-Monotonic Scheduling
 - Earliest-Deadline-First Scheduling
 - Proportional Share Scheduling
 - Example: POSIX Real-Time Scheduling

Minimizing Latency



- **Soft real-time computing**: requires that critical processes receive priority over less fortunate ones
- **Hard real-time systems**: requires to complete a critical task within a **guaranteed** amount of time
- **Event-driven** VS. interrupt-driven
 - The system is typically waiting for an event in real time to occur.
 - Events may arise either in software or in hardware
- **Event latency**: the amount of time from an event occurring to when it is serviced

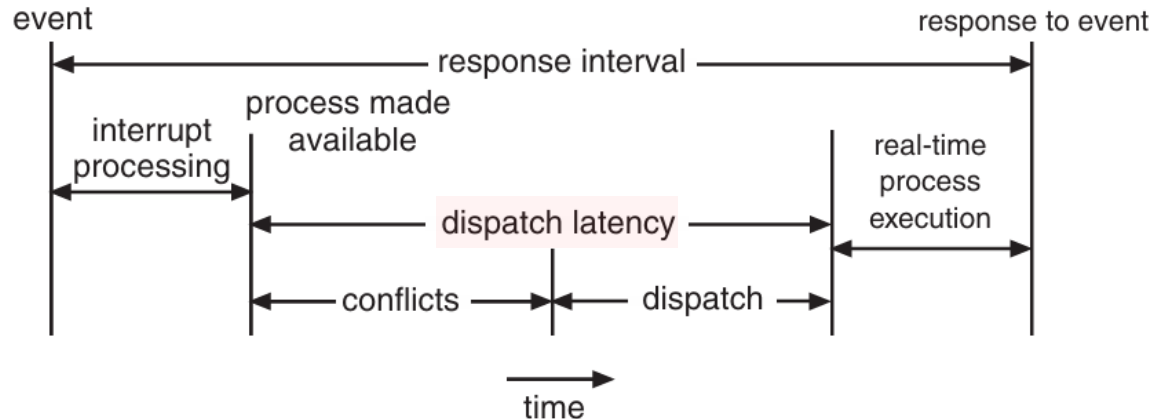


- Two types of latencies affect the performance of RTS
 - **Interrupt latency**: Time from the arrival of an interrupt at the CPU to the start of the ISR, should be minimized; For Hard RTS: should be bounded
 - 0. Interrupt may be disabled, the time should be very short and bounded.
When interrupt is enabled:
 1. Complete the current instruction
 2. Determine the interrupt type
 3. Save the state of the current process
 4. Start to execute the ISR
 - **Dispatch latency**

Minimizing Latency



- **Dispatch latency:** The amount of time required for the scheduling dispatcher to stop one process and start another




- Approaches for **short dispatch latency**
 - **Preemption**
 - preemption point (抢占点) in system calls with long period
 - preemptible kernel
 - **priority inversion**
 - priority-inheritance protocol
 - priority-ceiling protocol

Priority-Based Scheduling



- The scheduler for a RTOS must support a **priority based algorithm with preemption**.
 - Real-time processes: highest priority levels
 - Example: Windows [16-32]
 - Only guarantees **soft** real-time functionality
 - For **hard** real-time with deadline requirements, need **additional** scheduling feature.
- Usually, the processes are **periodic tasks**



Mostly, $t \leq d \leq p$

$$\text{Rate(速率)} = 1/p$$

- Schedulability: admission-control algorithm
 - Admit(=guarantee)? Or, reject?


RM scheduling



- Example1: $P_1: t_1 = 20, p_1 = d_1 = 50$; $P_2: t_2 = 35, p_2 = d_2 = 100$


CPU utilization

$$u_1 = t_1/p_1 = 20/50 = 0.4$$
$$u_2 = t_2/p_2 = 35/100 = 0.35$$
$$u_{total1} = u_1 + u_2 = 0.75$$




- Rate-Monotonic (RM) Scheduling

- using a static priority policy with preemption
- **Priority:** the shorter the period. the higher the priority



- Rate-monotonic scheduling is considered **optimal** in that
 - If a set of processes cannot be scheduled by this algorithm, it cannot be scheduled by any other algorithm that assigns **static** priorities.
- Example2: $P1: t_1 = 25, p_1 = d_1 = 50$; $P2: t_2 = 35, p_2 = d_2 = 80$



CPU utilization $u_{total2} = u_1 + u_2 = 25/50 + 35/80 = 0.94$

- Limitation:
 - CPU utilization is bounded, and it is not always possible to maximize CPU resources fully.
- The worst-case CPU utilization for scheduling N processes is

$$u_{wc, total} = N(2^{1/N} - 1)$$


- $N = 1$: $u_{wc, total} = 100\%$; $N = 2$: $u_{wc, total} \approx 83\%$
- $N \rightarrow \infty$: $u_{wc, total} \approx 69\%$

$$u_{total1} = 0.75$$


$$u_{total2} = 0.94$$

- Earliest-Deadline-First (EDF) Scheduling
 - Using dynamically priority according to deadline
 - May have to **adjust** the priorities when a new process became runnable to reflect its deadline
 - Priority: the earlier the deadline, the higher the priority
- Recall example2: $P1: t_1 = 25, p_1 = d_1 = 50$; $P2: t_2 = 35, p_2 = d_2 = 80$

✗ Scheduling3



Scheduling4





- EDF VS. RM
 - Periodic?
 - Constant t ?
 - Only deadline!
- The EDF scheduling is **theoretically optimal**
 - **theoretically**, guarantee & $u_{total} = 100\%$
 - But, **practically**, however, it is **impossible** due to the cost of context switching between processes and interrupt handling.

Proportional Share Scheduling



中国科学技术大学
University of Science and Technology of China

- **Proportional share scheduler:**
 - Total processor time: T shares
 - If an app can receive N shares, it will have N/T of the total processor time.
 - Example: $T = 100$, $N_1=50$, $N_2=15$, $N_3=20$
- **Admission-control policy:**
 - To guarantee an app receives its allocated share
 - If P_4 requested 30 shares, it will be denied.
 - $100-(50+15+20)=15 < 30$

POSIX Real-Time Scheduling



- The POSIX standard also provides extensions for real-time computing— POSIX.1b.
- POSIX defines two scheduling classes for real-time threads: (with equal priority)
 - `SCHED_FIFO`
 - `SCHED_RR`
- Related POSIX API:
 - `pthread_attr_getschedpolicy(pthread attr t *attr, int *policy)`
 - `pthread_attr_setschedpolicy(pthread attr t *attr, int policy)`



```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5

int main(int argc, char *argv[]) {
    int i, policy;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;

    /* get the default attributes */
    pthread_attr_init(&attr);

    /* get the current scheduling policy */
    if (pthread_attr_getschedpolicy(&attr, &policy) != 0)
        fprintf(stderr, "Unable to get policy.\n");
    else {
        if (policy == SCHED_OTHER)    printf("SCHED_OTHER\n");
        else if (policy == SCHED_RR)  printf("SCHED_RR\n");
        else if (policy == SCHED_FIFO) printf("SCHED_FIFO\n");
    }
}
```



```
/* set the scheduling policy- FIFO, RR, or OTHER */  
if (pthread_attr_setschedpolicy(&attr, SCHED_FIFO) != 0)  
    fprintf(stderr, "Unable to set policy.\n");  
  
/* create the threads */  
for (i = 0; i < NUM_THREADS; i++)  
    pthread_create(&tid[i], &attr, runner, NULL);  
  
/* now join on each thread */  
for (i = 0; i < NUM_THREADS; i++)  
    pthread_join(tid[i], NULL);  
}  
  
/* Each thread will begin control in this function */  
void *runner(void *param) {  
    /* do some work ... */  
    pthread_exit(0);  
}
```


- Real-time CPU scheduling
 - Performance: How to minimize the latencies
 - Event-driven & event latency
 - RTOS
 - Priority Scheduling with preemption
 - Short interrupt latency & short dispatch latency
 - Schedulability
 - Process modeling
 - Admission-control
 - Rate-Monotonic Scheduling
 - Earliest-Deadline-First Scheduling
 - Proportional Share Scheduling
 - Example: POSIX Real-Time Scheduling