



中国科学技术大学
University of Science and Technology of China

011174.01: Operating System 操作系统原理与设计

Chapter 4: Threads & Concurrency

陈香兰(xlanchen@ustc.edu.cn)

高能效智能计算实验室, CS, USTC @ 合肥

嵌入式系统实验室, CS, USTC @ 苏州



温馨提示:



为了您和他人的工作学习，
请在课堂上**关机或静音**。

不要在课堂上接打电话。

Chapter Objectives



中国科学技术大学
University of Science and Technology of China

- Chapter Objectives

1. **To introduce the notion of a thread** – a fundamental unit of CPU utilization that forms the basis of multithreaded computer system.
2. **To discuss the APIs for Pthreads, Win32, and JAVA thread libraries.**
3. Some threading issues.

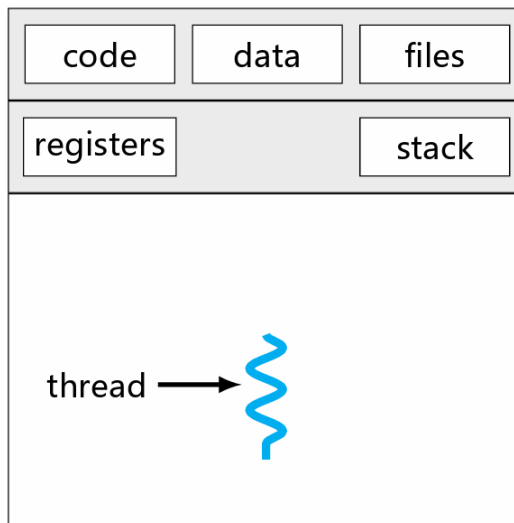
- **Overview**

- Single-threaded and multithreaded processes
 - Motivation
 - Benefits
- Multicore programming
 - Multithreading models
 - Thread Libraries
 - Threading Issues
 - OS examples for thread
 - Threading Scheduling

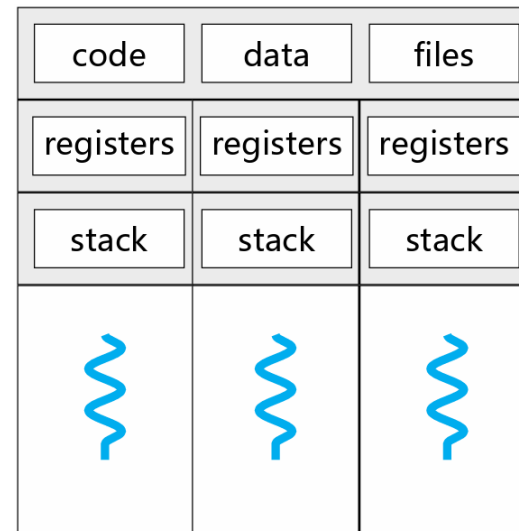
Thread concept overview



- A **thread** is a **basic unit of CPU utilization**;
 - It comprises a thread **ID**, a **program counter**, a **register set**, and a **stack**.
 - It shares with other threads belonging to the same process the code section, the data section, and other OS resources, such as open files, signals, etc.
- A **traditional process** has a single thread of control:
heavyweight process.



Single-threaded process



Multithreaded process

- On modern desktop PC, **many APPs are multithreaded.**
 - **A separate process with several threads**
 - Example 1: A web browser
 - One for displaying images or text;
 - another for retrieving data from network
 - Example 2: A word processor
 - One for displaying graphics;
 - Another for responding to keystrokes from the user;
 - And a third for performing spelling & grammar checking in the background
 - Example 3: RPC servers
 - For each message, a separate thread is used to service the message
 - concurrency↑



- Motivation

- In certain situations, a single application may be required to perform several similar tasks.
 - Example: a web server
- Allow a server to service several concurrent requests.
 - Example: an RPC server and Java' s RMI systems
- The OS itself needs to perform some specific tasks in kernel, such as managing devices or interrupt handling.
 - **PARTICULAR, many OS systems are now multithreaded.**
 - Example: Solaris, Linux

1. Responsiveness (响应性)

- Example: an interactive application such as web browser
 - while one thread loading an image, another allowing user interaction

2. Resource Sharing

- address space, memory, and other resources

3. Economy (经济, 指开销)

- Example: Solaris
 - creating a process is about 30 times slower than creating a thread;
 - context switching is about 5 times slower

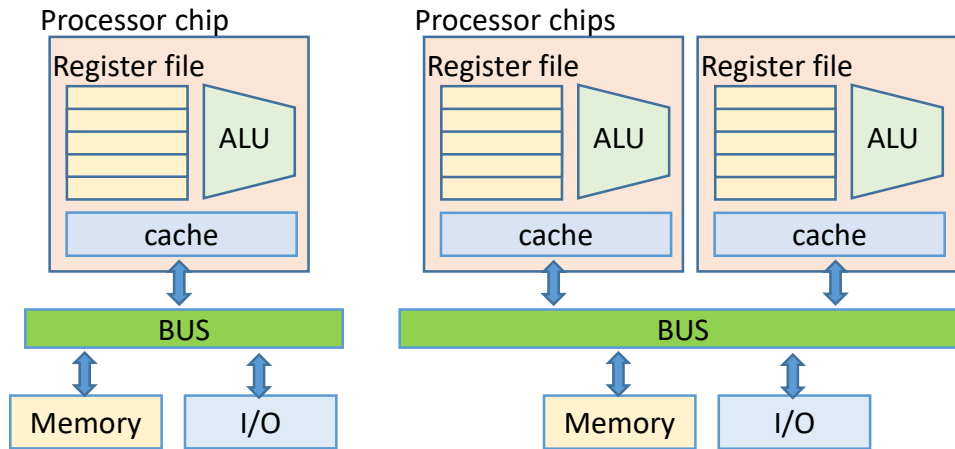
4. Scalability (可伸缩性)

- Utilization of MP Architectures;
- Parallelism and concurrency ↑

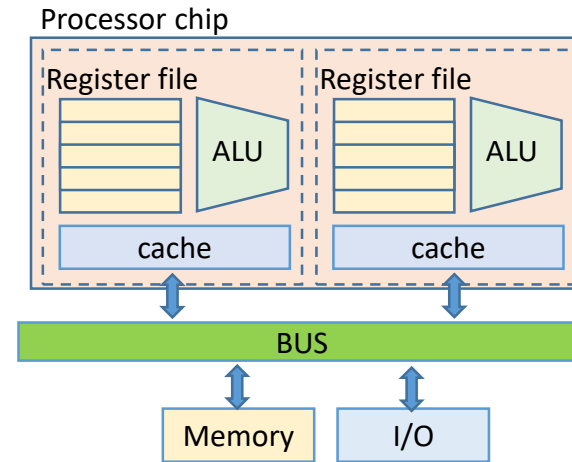


- Overview
- **Multicore programming**
 - Multicore and hyper-threading
 - Concurrency and parallelism
 - Programming challenges
 - Types of parallelism
- Multithreading models
- Thread Libraries
- Threading Issues
- OS examples for thread
- Threading Scheduling

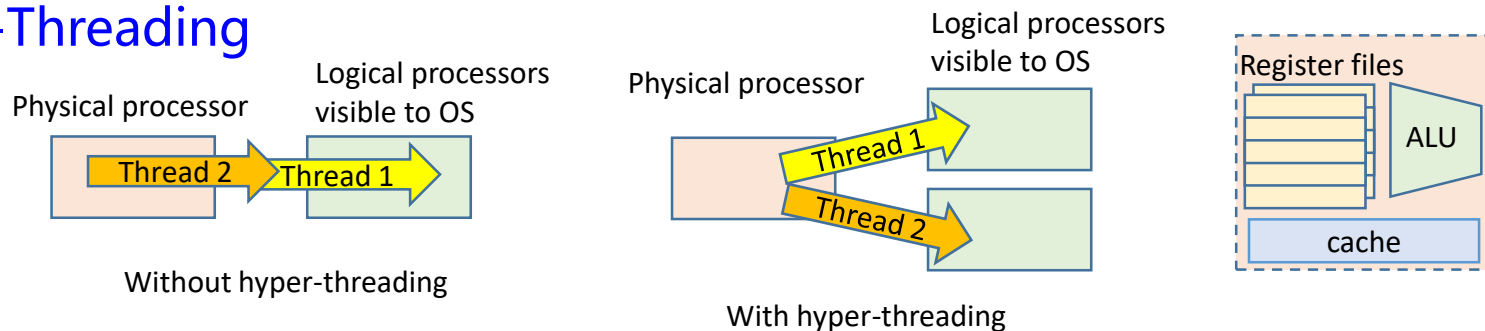
- Single-CPU VS. Multi-CPU



- Single-core VS. Multicore

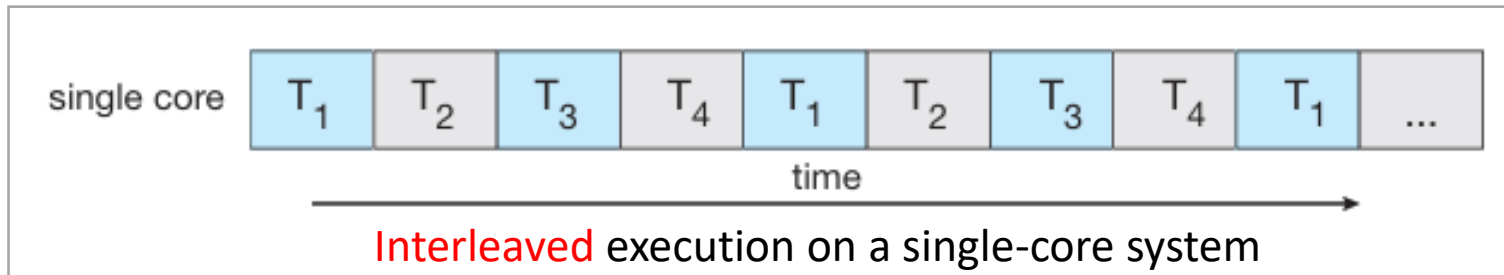


- Hyper-Threading




- Multithreaded programming** provides a **mechanism** for **more efficient use** of these multiple computing resources and **improved concurrency**.

- Concurrency VS parallelism




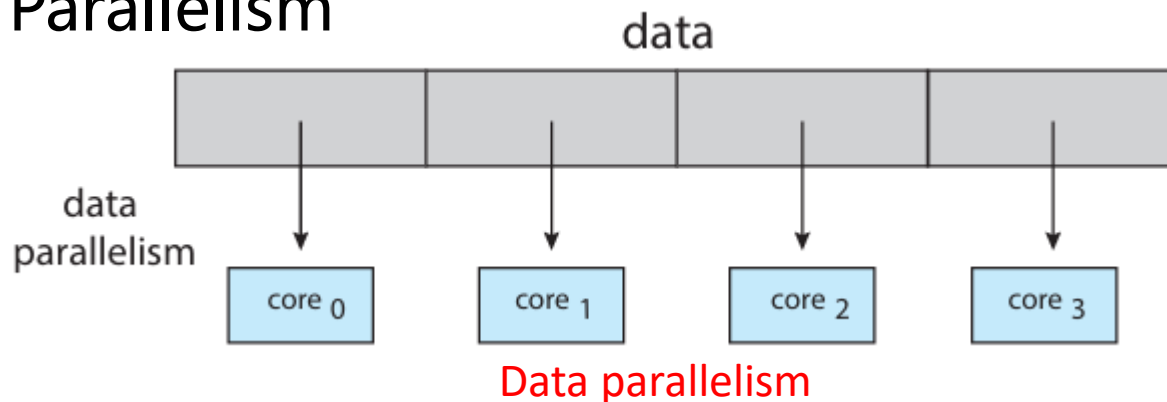
concurrency without parallelism



concurrency with parallelism

- Challenges in making better use of the multiple computing resources:
 - **OS designers:** What scheduling algorithms?
 - **APP programmers:** How to modify **existing programs** or to design **new programs** to be **multithreaded**?
 - **Identifying tasks:** how to **divide APP** into **separate, concurrent tasks** that can run in **parallel**?
 - **Balance:** how to ensure that the tasks perform **equal work of equal value**? Is each task **worth the cost** using a separate execution resource?
 - **Data splitting:** How to **divide the data** accessed and manipulated by the tasks to run on separate execution resources?
 - **Data dependency:** Any data dependencies among the tasks? How to examine? How to ensure the synchronized execution of the tasks?
 - **Testing & debugging:** How?

- Types of Parallelism




Data and task parallelism are not mutually exclusive, and an **application may** in fact **use a hybrid** of these two strategies.



- Overview
- Multicore programming
- **Multithreading models**
 - User threads VS kernel threads
 - Multithreading models
- Thread Libraries
- Threading Issues
- OS examples for thread

- Two methods to **support** threads
 - At the **user level**: **User threads**
 - Supported **above** the kernel; managed **by user-level threads library without** kernel support; Kernel may be multithreaded or not
 - Three primary thread libraries:
 1. POSIX Pthreads
 2. Win32 threads
 3. Java threads
 - By **kernel**: **Kernel threads**
 - Supported and managed **directly by the OS**.
 - Example: Virtually all contemporary OSes including
 - Windows, Linux, and macOS, ...




- The relationship between user threads and kernel threads

1. Many-to-One [n:1]
2. One-to-One [1:1]
3. Many-to-Many [n:m]


1. Many-to-One [n:1]

- Many user-level threads mapped to one kernel thread
- Examples:
 - Solaris Green Threads (adopted in early versions of Java)
 - GNU Portable Threads
- Disadvantages
 - the entire process will be **blocked** if a thread makes a blocking system call
 - unable to run in parallel** on systems with multiple computing resources
- Very few used




2. One-to-One [1:1]

- Each user-level thread maps to one kernel thread
- More concurrency, but not economy, too many kernel threads may burden the performance of a system.
- Examples
 - Windows NT/XP/2000
 - Linux
 - Solaris 9 and later




3. Many-to-Many [n:m]

- Allows many user level threads to be mapped to a smaller or equal number of kernel threads
- The number of kernel threads may be specific to either a particular APP or a particular machine.
- Examples
 - Solaris prior to version 9
 - Windows NT/2000 with the ThreadFiber package



4. Two-level Model, a popular variation on n:m model

- Similar to n:m, except that it allows a user thread to be bound to a kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier



- Overview
- Multicore programming
- Multithreading models
- **Thread Libraries**
 - Three main thread libraries
- Threading Issues
- OS examples for thread
- Threading Scheduling

Thread Libraries



- A thread library provides an API for creating and managing threads.
- Two primary ways
 1. To provide a library entirely in user space with no kernel support
 2. To implement a kernel-level library supported directly by the OS

| Library | Code & data | API | Invoking method inside API |
|--------------|------------------------|------------|----------------------------|
| User-level | Entirely in user space | User space | A local function call |
| Kernel-level | Kernel space | User space | System call |

- **Two general strategies**
 - **Asynchronous threading**
 - threads are independent with little data sharing
 - **Synchronous threading or fork-join strategy**
 - parent waits until all children have joined
 - involves significant data sharing among threads



- Three main thread libraries
 1. POSIX Pthreads
 2. Win32 threads
 3. Java threads
- Example: a multithreaded program
 - Summation:

$$sum = \sum_{i=0}^N i$$

1 Pthreads



- Pthreads
 - A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
 - API specifies behavior of the thread library, implementation is up to development of the library
 - Common in UNIX OSes (Solaris, Linux, Mac OS X)

Multithreaded C program using the Pthreads API



```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */

/* The thread will begin control in this function */
void *runner(void *param) {
    int i, upper = atoi(param);
    sum = 0;

    if (upper > 0) {
        for (i = 1; i <= upper; i++)
            sum += i;
    }
    pthread_exit(0);
}

int main(int argc, char *argv[]) {
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of attributes for the thread */
```


Multithreaded C program using the Pthreads API



```
if (argc != 2) {  
    fprintf(stderr, "usage: a.out \n" );  
    return-1;  
}  
  
if (atoi(argv[1]) < 0) {  
    fprintf(stderr, "Argument %d must be non-negative\n" ,atoi(argv[1]));  
    return-1;  
}  
  
pthread_attr_init(&attr); /* get the default attributes */  
pthread_create(&tid, &attr, runner, argv[1]); /* create the thread */  
pthread_join(tid, NULL); /* now wait for the thread to exit */  
  
printf( "sum = %d\n" , sum);  
}
```

NAME

pthread_attr_init, pthread_attr_destroy - initialise and destroy threads attribute object

SYNOPSIS

```
#include <pthread.h>
int pthread_attr_init(pthread_attr_t *attr);
int pthread_attr_destroy(pthread_attr_t *attr);
```

DESCRIPTION

The function pthread_attr_init() initialises a thread attributes object attr with the default value for all of the individual attributes used by a given implementation.

...

The pthread_attr_destroy() function is used to destroy a thread attributes object.

RETURN VALUE

Upon successful completion, both return a value of 0.
Otherwise, an error number is returned to indicate the error.

...

pthread_create()



中国科学技术大学
University of Science and Technology of China

NAME

pthread_create - thread creation

SYNOPSIS

```
#include <pthread.h>
```

```
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,  
                  void *(*start_routine)(void*), void *arg);
```

DESCRIPTION

The pthread_create() function is used to create a new thread, with attributes specified by attr, within a process. . . . Upon successful completion, pthread_create() stores the ID of the created thread in the location referenced by thread.

The thread is created executing start_routine with arg as its sole argument. . . .

. . .

If pthread_create() fails, no new thread is created and the contents of the location referenced by thread are undefined.

RETURN VALUE

If successful, the pthread_create() function returns zero.

Otherwise, an error number is returned to indicate the error.

NAME

pthread_join - wait for thread termination

SYNOPSIS

```
#include <pthread.h>
```

```
int pthread_join(pthread_t thread, void **value_ptr);
```

DESCRIPTION

The pthread_join() function suspends execution of the calling thread until the target thread terminates, unless the target thread has already terminated. ... The results of multiple simultaneous calls to pthread_join() specifying the same target thread are undefined. ...

RETURN VALUE

If successful, the pthread_join() function returns zero.
Otherwise, an error number is returned to indicate the error.
...

NAME

pthread_exit - thread termination

SYNOPSIS

```
#include <pthread.h>  
void pthread_exit(void *value_ptr);
```

DESCRIPTION

The pthread_exit() function terminates the calling thread and makes the value value_ptr available to any successful join with the terminating thread. ...

...

RETURN VALUE

The pthread_exit() function cannot return to its caller.

2 Win32 Threads Example



- Similar to the Pthreads technique.
- Multithreaded C program using the Win32 API

```
#include <stdio.h>
#include <windows.h>
DWORD Sum; /* data is shared by the thread(s) */

/* the thread runs in this separate function */
DWORD WINAPI Summation(PVOID Param){
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[]){
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;
    // do some basic error checking
```

2 Win32 Threads Example



```
if (argc != 2){
    fprintf(stderr, " An integer parameter is required\n" ); return-1;
}

Param = atoi(argv[1]);
if (Param < 0){
    fprintf(stderr, "an integer >= 0 is required \n" ); return-1;
}

// create the thread
ThreadHandle = CreateThread(NULL, //default security attribute
                            0, //default stack size
                            Summation, //thread function
                            &Param, //parameter to thread function
                            0, //default creation flags
                            &ThreadId);

if (ThreadHandle != NULL){
    WaitForSingleObject(ThreadHandle, INFINITE);
    CloseHandle(ThreadHandle);

    printf( "sum = %d\n" ,Sum);
}
}
```

3 Java Threads



- Java Threads
 - **Threads are the fundamental model** for program execution.
 - Java threads may be created by:
 - **Extending Thread class**
to create a new class that is derived from the Thread class and override its run() method.
 - **Implementing the Runnable interface**

Java Thread Example



```
class Sum {  
    private int sum;  
  
    public int get() {  
        return sum;  
    }  
  
    public void set(int sum) {  
        this.sum = sum;  
    }  
}
```

```
class Summation implements Runnable{  
    private int upper;  
    private Sum sumValue;  
  
    public Summation(int upper, Sum sumValue) {  
        if (upper < 0) throw new  
        IllegalArgumentException();  
        this.upper = upper;  
        this.sumValue = sumValue;  
    }  
  
    public void run() {  
        int sum = 0;  
        for (int i = 0; i <= upper; i++)  
            sum += i;  
        sumValue.set(sum);  
    }  
}
```

Java Thread Example



```
public class Driver {  
    public static void main(String[] args) {  
        if (args.length != 1) {  
            System.err.println( "Usage Driver <integer>" );  
            System.exit(0);  
        }  
  
        Sum sumObject = new Sum();  
        int upper = Integer.parseInt(args[0]);  
        Thread worker = new Thread(new Summation(upper, sumObject));  
        worker.start();  
        try {  
            worker.join();  
        } catch (InterruptedException ie) { }  
  
        System.out.println( "The sum of" + upper + " is " + sumObject.get());  
    }  
}
```

- Overview
- Multicore programming
- Multithreading models
- Thread Libraries
- **Threading Issues**
 - To discuss some of the issues to consider in designing multi threaded programs
- OS examples for thread
- Threading Scheduling

Threading Issues



中国科学技术大学
University of Science and Technology of China

- **Implicit threading**: (VS. explicit threading)
 - Developing multithreaded application is **hard, error-prone and time-consuming**
 - How to **transfer** the creation and management of threading from APPs developers **to compilers and run-time LIBs**.
 - Example:
 - Fork-join model; OpenMP; Grand Central Dispatch (大中央调度, GCD); Intel Thread Building Blocks (Intel的线程构建模块, TBB); ...
- **Thread Pools**:
 - At startup, create a number of threads in a pool where they await work
 - Advantages:
 - Usually slightly **faster** to service a request with an existing thread than create a new thread
 - Allows the number of threads in the APPs to **be bound to the size of the pool**
 - Separates the task from the mechanics of creating the task, **flexibility**↑



- **Semantics of fork() and exec() system calls**
 - Does fork() duplicate only the calling thread or all threads?
 - Some UNIX system have chosen to have two versions
 - Which one version to use? Depend on the APP.
- **Signal Handling** (ignore, or by yourself)
- **Thread cancellation** (线程撤销)
 - Terminating a thread before it has finished
 - **Two general approaches:**
 - **Asynchronous cancellation** terminates the target thread immediately
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled



- **Thread-Local Storage**

- Allows each thread to have its own copy of certain data
- similar to static data, but unique to each thread

- **Scheduler Activations** (调度器激活)

- Both n:m and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- Scheduler activations provide upcalls- a communication mechanism from the kernel to the thread library
- This communication allows an application to maintain the correct number kernel threads

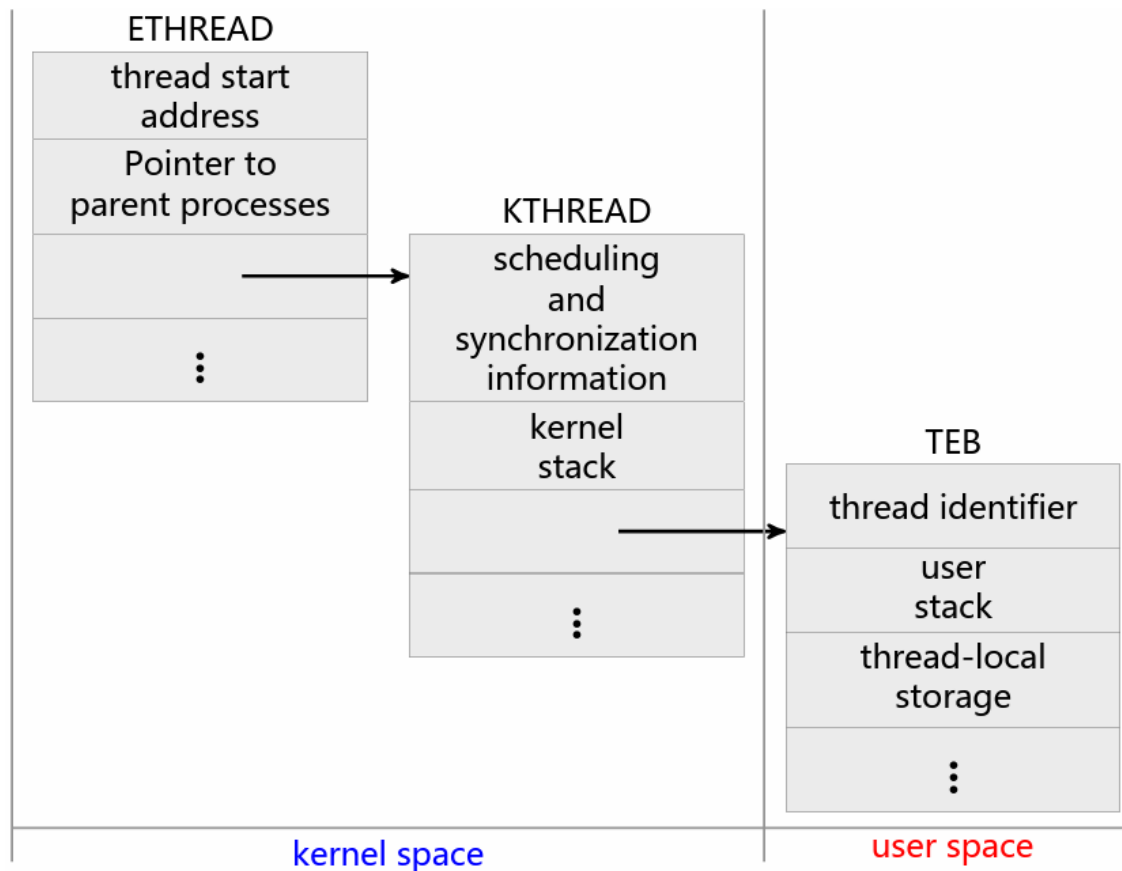


- Overview
- Multicore programming
- Multithreading models
- Thread Libraries
- Threading Issues
- **OS examples for thread**
 - Windows Threads
 - Linux Threads
 - Java Threads
- Threading Scheduling

Windows Threads




- The primary data structures of a thread include **ETHREAD** (executive thread block), **KTHREAD**(kernel thread block) ; **TEB**(thread environment block)



- Linux refers to them as tasks rather than threads
- Thread creation is done through **clone()** system call
- clone() allows **a child task to share the address space of the parent task (process)**
- clone() VS. fork()

| Flag | Meaning |
|---------------|-----------------------------------|
| CLONE_FS | File-system information is shared |
| CLONE_VM | The same memory space is shared |
| CLONE_SIGHAND | Signal handlers are shared |
| CLONE_FILES | The set of open files is shared |

- Java在**语言级**提供线程创建和管理支持功能
 - Java threads are managed by the JVM, not user-level library or kernel
- Java threads may be created by:
 - Extending Thread class
 - Implementing the Runnable interface Java



thead States



- Overview
- Multicore programming
- Multithreading models
- Thread Libraries
- Threading Issues
- OS examples for thread
- Threading Scheduling



- **user-level thread VS. kernel-level thread (or LWP)**
- **Local Scheduling**– How the **threads library** decides which thread to put onto an available LWP
 - many-to-one, many-to-many models
 - process-contention scope, PCS
- **Global Scheduling**– How the **kernel** decides which kernel thread to run next
 - many-to-one, many-to-many & one-to-one models
 - system-contention scope, SCS

- POSIX Pthread API **allows specifying either PCS** or SCS during thread creation
 - PTHREAD_SCOPE_PROCESS, many-to-many
 - PTHREAD_SCOPE_SYSTEM, one-to-one
 - create and **bind** an LWP for each user-level thread
- example

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5

int main(int argc, char *argv[]) {
    int i;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    pthread_attr_init(&attr); /* get the default attributes */

    /* set the scheduling algorithm to PROCESS or SYSTEM */
    pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);
```

Pthread Scheduling API



中国科学技术大学
University of Science and Technology of China

```
/* set the scheduling policy- FIFO, RT, or OTHER */
pthread_attr_setschedpolicy(&attr, SCHED_OTHER);

for (i = 0; i < NUM_THREADS; i++) /* create the threads */
    pthread_create(&tid[i], &attr, runner, NULL);

for (i = 0; i < NUM_THREADS; i++) /* now join on each thread */
    pthread_join(tid[i], NULL);
}

/* Each thread will begin control in this function */
void *runner(void *param) {
    printf( "I am a thread\n" );

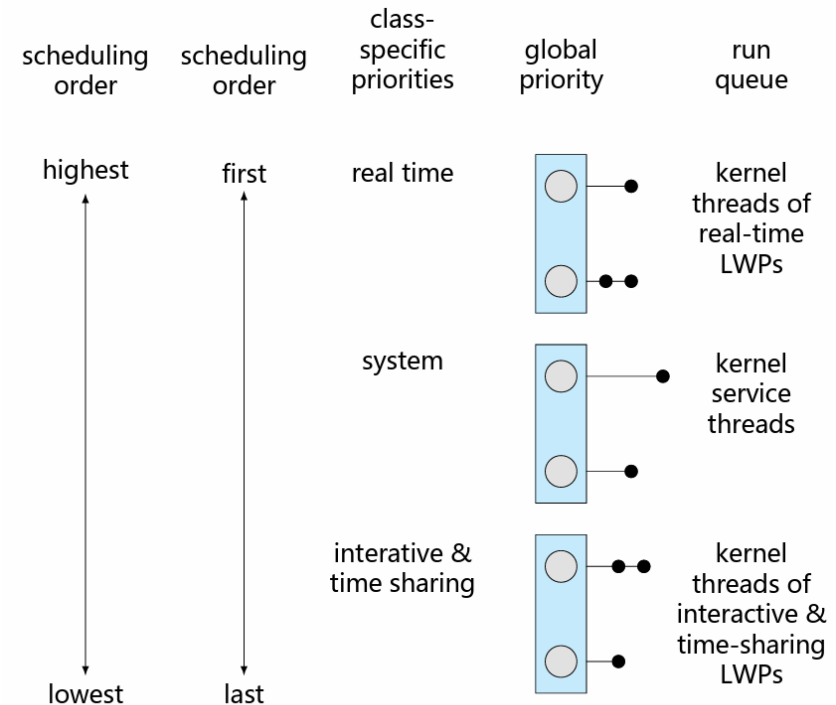
    pthread_exit(0);
}
```

Solaris scheduling



- Solaris: **priority-based** thread scheduling
- **4 classes of scheduling**, in order of priority. Within each class there are **different priorities** and **different scheduling algorithms**.

- **Real time**
- **System** (do not change the priority)
- **Time sharing** (default, with a multilevel feedback queue)
- **Interactive**, the same as time sharing, but higher priority



Solaris scheduling



| priority | time quantum | time quantum expired | return from sleep |
|----------|--------------|----------------------|-------------------|
| 0 | 200 | 0 | 50 |
| 5 | 200 | 0 | 50 |
| 10 | 160 | 0 | 51 |
| 15 | 160 | 5 | 51 |
| 20 | 120 | 10 | 52 |
| 25 | 120 | 15 | 52 |
| 30 | 80 | 20 | 53 |
| 35 | 80 | 25 | 54 |
| 40 | 40 | 30 | 55 |
| 45 | 40 | 35 | 56 |
| 50 | 40 | 40 | 58 |
| 55 | 40 | 45 | 58 |
| 59 | 20 | 49 | 59 |

Solaris Dispatch Table

Windows XP scheduling



- Dispatcher: **priority-based**, **preemptive** scheduling algorithm uses a **32-level priority** scheme to determine the order of thread execution
 - 0: idle thread
 - 1~15: variable classes priorities
 - 16~31: real-time class
 - One queue for each priority

| | real— time | high | above normal | normal | below normal | idle priority |
|---------------|---------------|------|-----------------|--------|-----------------|------------------|
| time-critical | 31 | 15 | 15 | 15 | 15 | 15 |
| highest | 26 | 15 | 12 | 10 | 8 | 6 |
| above normal | 25 | 14 | 11 | 9 | 7 | 5 |
| normal | 24 | 13 | 10 | 8 | 6 | 4 |
| below normal | 23 | 12 | 9 | 7 | 5 | 3 |
| lowest | 22 | 11 | 8 | 6 | 4 | 2 |
| idle | 16 | 1 | 1 | 1 | 1 | 1 |

Windows XP Priorities (**policy classes**, **relative priority**)

- Overview
- Multicore programming
- Multithreading models
- Thread Libraries
- Threading Issues
- OS examples for thread
- Threading Scheduling

Q & A