

0117401: Operating System

操作系统原理与设计

Chapter 11: File system implementation(文件系统实现)

陈香兰

xlanchen@ustc.edu.cn

<http://staff.ustc.edu.cn/~xlanchen>

Computer Application Laboratory, CS, USTC @ Hefei
Embedded System Laboratory, CS, USTC @ Suzhou

May 15, 2024

温馨提示：



为了您和他人的工作学习，
请在课堂上**关机或静音**。

不要在课堂上接打电话。

提纲

File-System Structure

FS Implementation

Directory Implementation

Allocation Methods (分配方法)

Free-Space Management

Efficiency (空间) and Performance (时间)

Recovery

Log Structured File Systems

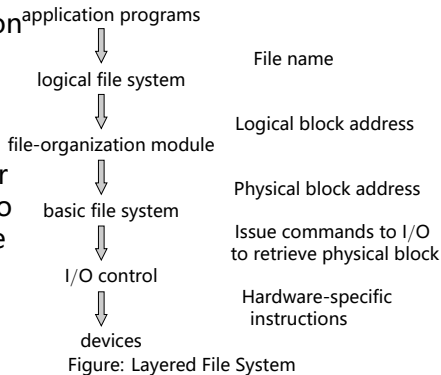
小结

Outline

File-System Structure

File-System Structure

- ▶ File structure
 - ▶ Logical storage unit
 - ▶ Collection of related information
- ▶ FS resides on secondary storage (disks)
- ▶ FS organization
 - ▶ How FS should look to the user
 - ▶ How to map the logical FS onto the physical secondary-storage devices
- ▶ FS organized into layers



Outline

FS Implementation

FS Implementation

- ▶ Structures and operations used to implement file system operation, OS- & FS-dependment
 1. On-disk structures
 2. In-memory structures

FS Implementation

1. On-disk structures

1.1 Boot control block

- ▶ To boot an OS from the partition (volume)
- ▶ If empty, no OS is contained on the partition

1.2 Volume control block

1.3 Directory structure

1.4 Per-file FCB

file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks

Figure: A typical file control block

FS Implementation

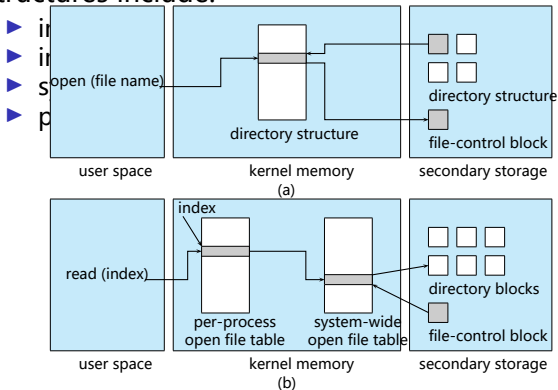
2. **In-memory information**: For both FS management and performance improvement via caching

- ▶ Data are loaded at mount time and discarded at dismount
- ▶ Structures include:
 - ▶ in-memory mount table;
 - ▶ in-memory directory-structure cache
 - ▶ system-wide open-file table;
 - ▶ per-process open-file table

FS Implementation

2. **In-memory information:** For both FS management and performance improvement via caching

- ▶ Data are loaded at mount time and discarded at dismount
- ▶ Structures include:




Partitions and mounting

- ▶ Partition (分区)
 - ▶ Raw (E.g. UNIX swap space & some database) VS. cooked
 - ▶ **Boot** information, with its own format
 - ▶ Boot image
 - ▶ Boot loader unstanding multiple FSES & OSes
Dual-boot
- ▶ **Root partition is mounted at boot time**
- ▶ **Others** can be automatically mounted at boot or manually mounted later

Virtual File Systems (虚拟文件系统)

- ▶ **Virtual File Systems (VFS, 虚拟文件系统)** provide an object-oriented way of implementing file systems.
- ▶ VFS allows the same system call interface (the API) to be used for different types of file systems.
- ▶ The API is to the VFS interface, rather than any specific type of file system.



Schematic View of Virtual File System

Outline

Directory Implementation

Directory Implementation

1. **Linear list** of file names with pointer to the data blocks.
 - ▶ Simple to program
 - ▶ Time-consuming to execute
2. **Hash Table**— linear list with hash data structure.
 - ▶ Decreases directory search time
 - ▶ **Collisions** – situations where two file names hash to the same location
 - ▶ Fixed & variable size or chained-overflow hash table

Outline

Allocation Methods (分配方法)


Allocation Methods (分配方法)

- ▶ An allocation method refers to **how disk blocks are allocated** for files so that disk space is **utilized effectively** & files can be **accessed quickly**
 1. Contiguous allocation (连续分配)
 2. Linked allocation (链接分配)
 3. Indexed allocation (索引分配)
 4. Combined (组合方式)

1. Contiguous Allocation (连续分配) I

- ▶ Each file occupies a **set of contiguous blocks on the disk**
- ▶ Simple – directory entry only need
 - ▶ **starting location (block #)**
 - ▶ **& length (number of blocks)**
- ▶ Mapping from logical to physical
 - $\text{LogicalAddress}/512 = Q \dots R$**
 - Block to be accessed = $Q + \text{starting address}$**
 - Displacement into block = R**

1. Contiguous Allocation (连续分配) II



directory		
file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2

1. Contiguous Allocation (连续分配) III

- ▶ Advantages:

- ▶ Support **both random & sequential** access

- ▶ Start block: b ;

- Logical block number: i

- \Rightarrow physical block number: $b + i$

- ▶ **Fast** access speed, because of short head movement

- ▶ Disadvantages:

- ▶ **External fragmentation**

- ▶ **Wasteful of space** (dynamic storage-allocation problem).

- ▶ **Files cannot grow,**

- or File size must be known in advance.

- \Rightarrow **Internal fragmentation**

Extent-Based Systems

- ▶ Many newer file systems (l.e. Veritas File System) use a **modified contiguous allocation scheme**
- ▶ Extent-based file systems allocate disk blocks in extents
- ▶ An extent is a contiguous block of disks
 - ▶ Extents are allocated for file allocation
 - ▶ A file consists of one or more extents.

2. Linked Allocation (链接分配)

- ▶ Each file is a linked list of disk blocks: blocks may be scattered anywhere on the disk.
- ▶ Two types
 1. **Implicit (隐式链接)**
 2. **Explicit (显式链接)**

2. Linked Allocation (链接分配)


1. Implicit (隐式链接)

- ▶ **Directory** contains a pointer to the first block & last block of the file.
- ▶ **Each block** contains a pointer to the next block.

a block =

pointer

- ▶ **Allocate as needed, link together**
 - ▶ Simple – need only starting address
 - ▶ Free-space management system – **no waste of space**



2. Linked Allocation (链接分配)

1. Implicit (隐式链接)

► Disadvantage:


► No random access

► Link pointers need disk sapce

E.g.: 512 per block, 4 per pointer $\Rightarrow 0.78\%$

Solution: clusters

\Rightarrow disk throughput \uparrow
But internal fragmentation \uparrow



2. Linked Allocation (链接分配)

1. Implicit (隐式链接)

► Mapping:

Suppose

- 1.1 block size=512B,
- 1.2 block pointer size=1B, using the first byte of a block
- 1.3 Logical addr in the file to be accessed= A


we have

- 1.1 Data size for each block = $512 - 1 = 511$
- 1.2 $A/511 = Q \dots R$

then

- 1.1 Block to be accessed is the Q^{th} block in the linked chain of blocks representing the file.
- 1.2 Displacement into block = $R + 1$


► How to reduce searching time?



2. Linked Allocation (链接分配)

2. **Explicit linked allocation:** **File Allocation table, FAT**

Disk-space allocation used by MS-DOS and OS/2




► **A section of disk at the beginning of each partition is set aside to contain the FAT**

- Each disk block one entry
- The entry contains
 - (1) **the index of the next block** in the file
 - (2) **end-of-file**, for the last block entry
 - (3) **0**, for unused block

► **Directory entry** contains the first block number


no. of disk blocks -1



2. Linked Allocation (链接分配)

2. **Explicit linked allocation:** **File Allocation table, FAT**

Disk-space allocation used by MS-DOS and OS/2




- ▶ Now **support random access**, but still not very efficient
- ▶ May result in a **significant disk head seeks**.

Solution: Cached FAT

2. Linked Allocation (链接分配)

2. **Explicit linked allocation:** **File Allocation table, FAT**

Disk-space allocation used by MS-DOS and OS/2



► **How to compute FAT size?**

Suppose

2.1 Disk space = 80 GB

2.2 Block size = 4 KB

Then

2.1 Total block number = $80 \times 2^{30} / 2^{12} = 5 \times 2^{22}$

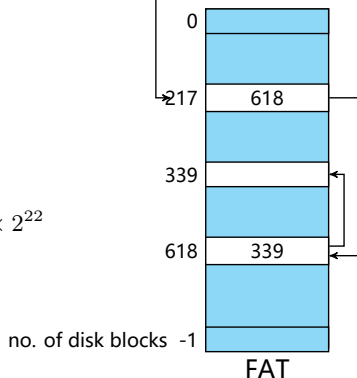
2.2 $4 \times 2^{22} = 2^{24} < 5 \times 2^{22} < 8 \times 2^{22} = 2^{25}$

► **Length of each FAT entry?**

(25bits? 28bits? 32bits?)

► **Length of FAT?**

$(5 \times 2^{22} \times 4B = 80MB = 80GB / 2^{10})$




3. Indexed Allocation (索引分配)

- ▶ **Indexed Allocation (索引分配):**

Brings all pointers together into one location – **the index block**.

- ▶ **Each file** has its own index block
- ▶ **Directory entry** contains the index block address
- ▶ **Each index block:** An array of pointers (an index table)

Logical block number i
= the i^{th} pointer



3. Indexed Allocation (索引分配)

- ▶ **Indexed Allocation (索引分配):**


Brings all pointers together into one location – **the index block**.

- ▶ **Advantage:**

- ▶ **Random** access
- ▶ Dynamic access **without external fragmentation**

- ▶ **Disadvantage:**

- ▶ have **overhead** of index block.
- ▶ File **size limitation**, since one index block can contain limited pointers



3. Indexed Allocation (索引分配)

► Indexed Allocation (索引分配):

Brings all pointers together into one location – **the index block**.

► Mapping from logical to physical

Suppose

(1) Block size = 1KB

(2) Index size = 4B

Then for logical address LA, we have


$$LA/512 = Q \dots R$$

(3) Q = the index of the pointer

(4) R = displacement into block

We also have **Max file size**

$$= 2^{10}/4 \times 1KB = 256KB$$




3. Indexed Allocation (索引分配)

► **Indexed Allocation (索引分配):**

Brings all pointers together into one location – **the index block**.

► **How to support a file of unbounded length?**

1. **linked scheme**
2. **multi-level index scheme**



3. Indexed Allocation (索引分配)

1. Linked scheme

- ▶ **Link blocks of index table (no limit on size).**
- ▶ **Mapping**

Suppose

(1) Block size=1KB

(2) Index or link pointer size = 4B

Then

$$LA / (1KB \times (1K/4 - 1)) = Q_1 \dots R_1$$

(3) Q_1 = block of index table

(4) R_1 is used as follows:

$$R_1 / 1K = Q_2 \dots R_2$$

(5) Q_2 = index into block of index table

(6) R_2 = displacement into block of file:

3. Indexed Allocation (索引分配)

2. multi-level index scheme

Example: **Two-level index** (maximum file size is ?)

► We have

$$LA / (1K \times 1K / 4) = Q_1 \dots R_1$$


(1) Q_1 = index into outer-index

(2) R_1 is used as follows:

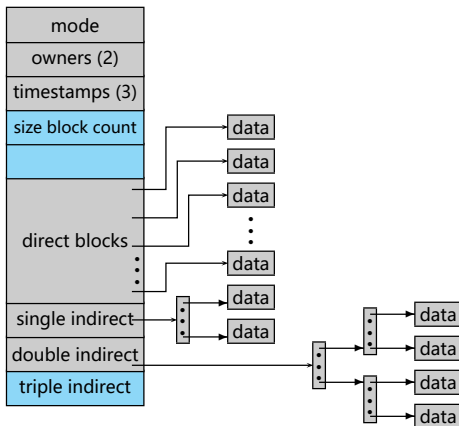
$$R_1 / 1KB = Q_2 \dots R_2$$

(3) Q_2 = displacement into block of index table

(4) R_2 = displacement into block of file



4. Combined Scheme (组合方式): UNIX (4K bytes per block) I



4. Combined Scheme (组合方式): UNIX (4K bytes per block) II

- ▶ if 4KB per block, and 4B per entry

$$\text{Direct blocks} = 10 \times 4\text{KB} = 40\text{KB}$$

$$\text{Number of entries per block} = 4\text{KB}/4\text{B} = 1\text{K}$$

$$\text{Single indirect} = 1\text{K} \times 4\text{KB} = 4\text{MB}$$

$$\text{Double indirect} = 1\text{K} \times 4\text{MB} = 4\text{GB}$$

$$\text{Triple indirect} = 1\text{K} \times 4\text{GB} = 4\text{TB}$$

Maximnm file size = ?

Outline

Free-Space Management

Free-Space Management

- ▶ **Disk Space**: limited
 - ▶ Free space management: **To keep track of free disk space**
 - ▶ **How?** Free-space list?
 - ▶ **Algorithms**
 1. **Bit vector**
 2. **Linked list**
 3. **Grouping (成组链接法)**
 4. **Counting**

Free-Space Management

1. Bit vector

- ▶ **Free-space list** is implemented as a **bit map** or **bit vector**
 - ▶ **1 bit for each block**
1=free;
0=allocated
 - ▶ Example:
a disk where blocks
2,3,4,5,8,9,10,11,12,13,17,18,25,26,27 are free and the
rest blocks are allocated. The bitmap would be

0011 1100 1111 1100 0110 0000 0111 0000 0...

- ▶ **Bit map length.**
For n blocks, if the base unit is word, and the size of
word is 16 bits, then

$$\text{bit map length} = (n + 15)/16$$

U16 bitMap[bitMapLength];

Free-Space Management

1. Bit vector

► How to find the first free block or n consecutive free blocks on the disk?

- Many computers supply **bit-manipulation instructions**
- To find the first free block:
Suppose: base unit = word (**16** bits) or other
 - (1) find **the first non-0 word**
 - (2) find **the first 1 bit** in the first non-0 word
- If first K words is 0, & $(K + 1)^{\text{th}}$ word > 0 ,
the first $(K + 1)^{\text{th}}$ word's first 1 bit has offset L ,
then

$$\text{first free block number } N = K \times 16 + L$$

Free-Space Management

1. Bit vector

- ▶ **Simple**

- ▶ Must be kept on disk

Bit map requires extra space,

Example:

block size = 2^{12} bytes

disk size = 2^{30} bytes (1 gigabyte)

$n = 2^{30} / 2^{12} = 2^{18}$ bits (or 32K bytes)

- ▶ **Solution: Clustering**

Free-Space Management


1. Bit vector

- ▶ **Efficient** to get the first free block or n consecutive free blocks, **if we can always store the vector in memory.**
 - ▶ But **copy in memory and disk may differ.**
E.g. **bit[i] = 1 in memory & bit[i] = 0 on disk**
 - ▶ **Solution:**
 - Set bit[i] = 1 in memory.
 - Allocate block[i]
 - Set bit[i] = 1 in disk
- ▶ Need to protect:
 - ▶ Pointer to free list
 - ▶ Bit map

Free-Space Management

2. Linked Free Space List on Disk

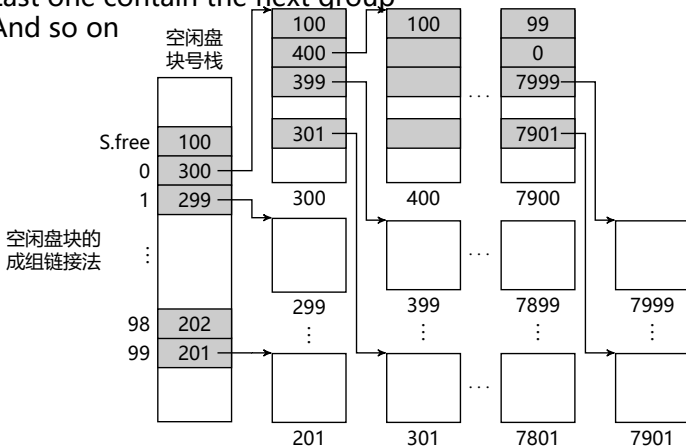
- ▶ **Link together all the free disk blocks**
 - ▶ **First** free block
 - ▶ **Next** pointer
- ▶ Not efficient
- ▶ **Cannot get contiguous space easily**
- ▶ No waste of space



Free-Space Management

3. **Grouping (成组链接法)**: To store the addresses of n free blocks (a group) in the first free block. E.g.: UNIX

- ▶ First $n-1$ group members are actually free
- ▶ Last one contain the next group
- ▶ And so on



Free-Space Management

4. Counting

- ▶ Assume:
Several contiguous blocks may be allocated or freed simultaneously
- ▶ Each = first free block number & a counter (number of free blocks)
- ▶ Shorter than linked list at most time, generally counter > 1

Outline

Efficiency (空间) and Performance (时间)

1 Efficiency (空间)


Efficiency in usage of disk space dependent on:

1. Disk allocation and directory algorithms
2. **Various approaches**
 - ▶ Inodes distribution
 - ▶ Variable cluster size
 - ▶ Types of data kept in file's directory entry
 - ▶ Large pointers provides larger file length, but cost more disk space

2 Performance (时间)

- ▶ **Performance:** other ways
 - ▶ **disk cache** - on disk controllers, large enough to store **entire tracks** at a time.
 - ▶ **buffer cache** – separate section of main memory for frequently used **blocks**
 - ▶ **page cache** - uses virtual memory techniques to **cache file data as pages** rather than as file-system-oriented blocks
 - ▶ **Synchronous writes** VS. **Asynchronous writes**
 - ▶ **free-behind and read-ahead** – techniques to optimize **sequential access**
 - ▶ improve PC performance by dedicating section of memory as virtual disk, or **RAM disk**

2 Performance (时间)



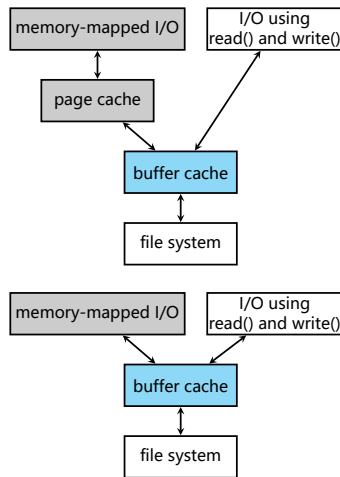
Unified Buffer Cache

► I/O Without a Unified Buffer Cache

- Memory-mapped I/O uses a page cache
- Routine I/O through the file system uses the buffer (disk) cache
- **Problem:** double caching

► I/O Using a Unified Buffer Cache

- A unified buffer cache uses the same page cache to cache both memory-mapped pages and ordinary file system I/O



Outline

Recovery

Recovery

- ▶ **Consistency checking** (一致性检查)
 - ▶ **compares** data in directory structure with data blocks on disk, and **tries to fix** inconsistencies
 - ▶ UNIX: fsck
 - ▶ MS-DOS: chkdsk
- ▶ **Backup & restore**
 - ▶ Use system programs to back up data from disk to another storage device (floppy disk, magnetic tape, other magnetic disk, optical)
 - ▶ Recover lost file or disk by restoring data from backup
 - ▶ A typical backup schedule may be:
 - Day1: **full backup**;
 - Day2: **incremental backup**;
 - ...
 - DayN: incremental backup. Then go back to Day1.

Outline

Log Structured File Systems

Log Structured File Systems

- ▶ **Log-based transaction-oriented** (or journaling, 日志) file systems record each update to the file system as a **transaction**
- ▶ All transactions are written to a log
 - ▶ A transaction is considered committed once it is written to the log
 - ▶ However, the file system may not yet be updated
- ▶ The transactions in the log are **asynchronously** written to the file system
 - ▶ When the file system is modified, the transaction is removed from the log
- ▶ If the file system crashes, all remaining transactions in the log must still be performed

Outline

小结

小结

File-System Structure

FS Implementation

Directory Implementation

Allocation Methods (分配方法)

Free-Space Management

Efficiency (空间) and Performance (时间)

Recovery

Log Structured File Systems

小结

Thank you! Any question?