

# 0117401: Operating System 操作系统原理与设计

## Chapter 13: IO Systems (IO管理)

陈香兰

[xlanchen@ustc.edu.cn](mailto:xlanchen@ustc.edu.cn)

<http://staff.ustc.edu.cn/~xlanchen>

Computer Application Laboratory, CS, USTC @ Hefei  
Embedded System Laboratory, CS, USTC @ Suzhou

May 29, 2024

# 温馨提示:



为了您和他人的工作学习,  
请在课堂上**关机或静音**。

**不要**在课堂上接打电话。

# 提纲

## 1 I/O Hardware and I/O control methods

- Polling (轮询方式)
- Interrupts (中断方式)
- Direct Memory Access (DMA方式)
- I/O hardware summary

## 2 Application I/O Interface

- Block and Character Devices
- Network Devices
- Clocks and Timers
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

## 3 Kernel I/O Subsystem

- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- Error Handling
- I/O Protection
- Kernel Data Structures

## 4 Transforming I/O Requests to Hardware Operations

## 5 Performance

## 6 小结

# Overview

- The role of OS in I/O is to manage and control I/O operations and I/O devices connected to the computer.
- **Challenge:** I/O devices **vary widely**.
- **HOW:** a combination of HW and SW techniques.

## Chapter Objectives

- To explore the structure of an OS' s I/O subsystem.
- To discuss the principles of I/O HW and its complexity.
- To provide details of the performance aspects of I/O HW and SW.

# Outline

- 1 I/O Hardware and I/O control methods
  - Polling (轮询方式)
  - Interrupts (中断方式)
  - Direct Memory Access (DMA方式)
  - I/O hardware summary

# I/O Hardware overview

- Incredible variety of I/O devices

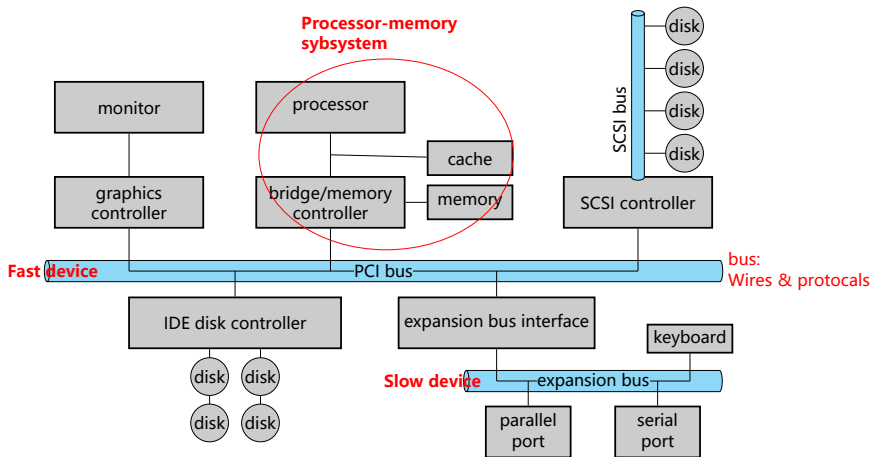


Figure: A typical PC bus structure.

# I/O Hardware overview

- Common concepts : **CPU→PORT→BUS→Controller**
  - ▶ **Port (端口)**: the connection point via which a device communicates with the machine.
  - ▶ **Bus (总线)**: daisy chain(菊花链) or shared direct access
    - ★ PCI (Peripheral Component Interconnect(外部器件互连) )
    - ★ SCSI (Small computer systems interface)
    - ★ Expansion bus
  - ▶ **Controller (控制器)**
    - ★ Simple: serial port
    - ★ Complex: bus controller (host adapter), device controller

# I/O Hardware overview

- How can the processor command controller?
  - ▶ **Controller** has one or more registers for data and control signals.
  - ▶ The processor communicates with the controller by reading and writing bit patterns in the registers.
- Two communication techniques:
  - 1 **Direct I/O instructions**
    - ★ Access the port address
    - ★ Each port typically contains of four registers, i.e., status, control, data-in and data-out.
    - ★ Instructions: In, out
  - 2 **Memory-mapped I/O**
    - ★ Example: 0xa0000 ~ 0xfffff are reserved to ISA graphics cards and BIOS routines
- ▶ Some systems use both techniques: PC as an example.



# I/O Hardware overview

- I/O address range

## Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000-00F	DMA controller
020-021	interrupt controller
040-043	timer
200-20F	game controller
2F8-2FF	serial port (secondary)
320-32F	hard-disk controller
378-37F	parallel port
3D0-3DF	graphics controller
3F0-3F7	diskette-drive controller
3F8-3FF	serial port (primary)

# I/O Control Methods

- 1 Polling (轮询方式)
- 2 Interrupts (中断方式)
- 3 DMA (DMA方式)
- 4 (在汤书上：还有通道的概念)

# Outline

- 1 I/O Hardware and I/O control methods
  - Polling (轮询方式)
  - Interrupts (中断方式)
  - Direct Memory Access (DMA方式)
  - I/O hardware summary

# Polling (轮询方式)

- Need **handshaking (握手)**
- State of device
  - 1 **command-ready**
    - ★ In command register
    - ★ 1: a command is available for the controller
  - 2 **busy**
    - ★ In status register
    - ★ 0: ready for the next command; 1: busy
  - 3 **Error**
    - ★ To indicate whether an I/O is ok.

# Polling (轮询方式)

- **Basic handshaking** notion for writing output
  - ① Host repeatedly **reads** the **busy** bit until it is 0
  - ② Host **sets write** bit in command register and writes a byte into data-out register
  - ③ Host **sets command-ready** bit
  - ④ When controller notices command-ready bit, it **sets busy** bit
  - ⑤ Controller gets write command and data, and works
  - ⑥ Controller **clears** command-ready bit, error bit and busy bit
- Step1: Busy-wait cycle to wait for I/O from device  
≡**busy-waiting**≡**polling**

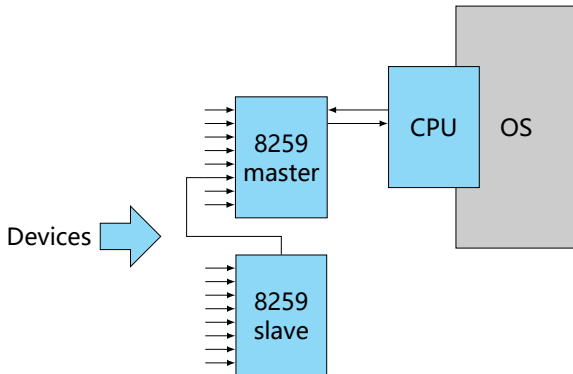
# Outline

## 1 I/O Hardware and I/O control methods

- Polling (轮询方式)
- **Interrupts (中断方式)**
- Direct Memory Access (DMA方式)
- I/O hardware summary

# Interrupts (中断方式)

- CPU Interrupt-request line triggered by I/O device
- Interrupt handler receives interrupts
- Basic interrupt scheme
  - ▶ Raise → Catch → Dispatch → Clear



# Interrupts (中断方式)

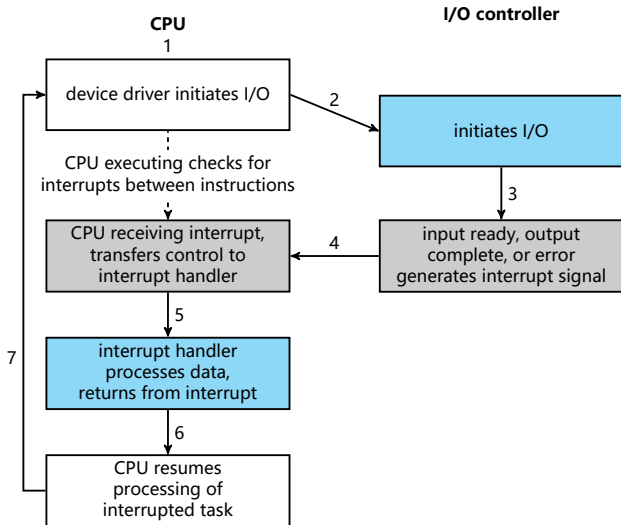


Figure: Interrupt-Driven I/O Cycle



# Interrupts (中断方式)

- More sophisticated interrupt-handling features:  
Most CPU have **two interrupt request lines**.
  - 1 **Nonmaskable**
  - 2 **Maskable** to ignore or delay some interrupts
- Efficient dispatching without polling the devices
  - ▶ **Interrupt vector**: to dispatch interrupt to correct handler
  - ▶ **Interrupt chaining**: to allow more device & more interrupt handlers
- Distinguish between high- and low-priority interrupts:
  - ▶ **Interrupt priority**: the handling of low-priority interrupts is deferred without masking, even preempted.
- Interrupt mechanism also used for **exceptions**

# Interrupts (中断方式)

- Example: Intel Pentium Processor Event-Vector Table

vector number	description	vector number	description
0	divide error	11	segment no present
1	debug exception	12	stack fault
2	null interrupt	13	general protection
3	breakpoint	14	page fault
4	INTO-detected overflow	15	(Intel reserved, do not use)
5	bound range exception	16	floating-point error
6	invalid opcode	17	alignment check
7	device not available	18	machine check
8	double fault	19-31	(Intel reserved, do not use)
9	coprocessor segment overrun (reserved)	32-255	maskable interrupts
10	invalid task state segment		

# Outline

- 1 I/O Hardware and I/O control methods
  - Polling (轮询方式)
  - Interrupts (中断方式)
  - **Direct Memory Access (DMA方式)**
  - I/O hardware summary

# Direct Memory Access (DMA方式)

- **Direct Memory Access (DMA方式):**

Used to avoid programmed I/O **for large data movement**, and bypasses CPU to **transfer data directly between I/O device and memory**

- Requires **DMA controller**

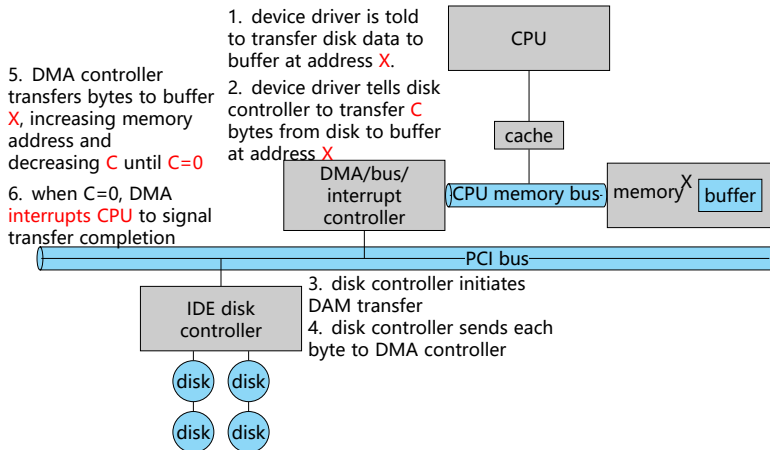
- ▶ the host prepares a **DMA command block** in memory
  - ★ a pointer to the source of a transfer
  - ★ a pointer to the destination of the transfer
  - ★ a count of the number of bytes to be transferred
- ▶ CPU writes the address of the DMA command block to DMA controller, and then goes on with other work.

# Direct Memory Access (DMA方式)

- Handshaking between DMA controller & device controller
  - ① Device controller **raises DMA-request** when one word is available
  - ② DMA controller **seizes memory bus**, places the desired address on memory-address wires, and raises DMA-acknowledge
  - ③ Device controller **transfers** the word to memory, and removes the DMA-request signal. Goto 1
  - ④ DMA controller **interrupts** the CPU.

# Direct Memory Access (DMA方式)

- Six Steps in a DMA transfer



- **Cycle stealing**: when DMA seizes the memory bus, CPU is momentarily prevented from accessing main memory

# Outline

- 1 I/O Hardware and I/O control methods
  - Polling (轮询方式)
  - Interrupts (中断方式)
  - Direct Memory Access (DMA方式)
  - I/O hardware summary

# I/O hardware summary

- A bus
- A controller
- An I/O port and its registers
- The handshaking relationship between the host and a device controller
- The execution of this handshaking in a pooling loop via interrupts
- the offloading of this work to a DMA controller for large transfer



# Outline

## 2 Application I/O Interface

- Block and Character Devices
- Network Devices
- Clocks and Timers
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

# I/O control challenges

- **Two challenges**

Applications  $\rightarrow$  OS  $\leftarrow$  Devices

- ▶ How can the OS give a convenient, uniform I/O interface to applications?
- ▶ How can the OS be designed such that new devices can be attached to the computer without the OS being rewritten?

# I/O control challenges

- **Device-driver Layer** hides differences among device controllers from the I/O subsystem of the kernel

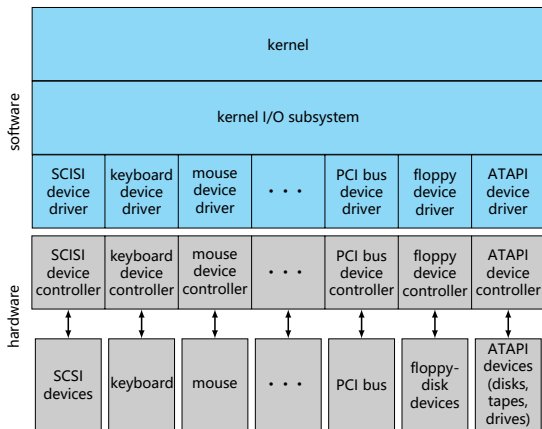


Figure: A Kernel I/O Structure

# Application I/O Interface

- **I/O system calls** encapsulate the behavior of devices in a few generic classes that hide HW differences from APPs.
- 设备独立性【汤】：应用程序与具体的物理设备无关。
- Devices vary in many dimensions
  - ▶ Character-stream or block
  - ▶ Sequential or random-access
  - ▶ Sharable or dedicated
  - ▶ Speed of operation
  - ▶ read-write, read only, or write only

# Characteristics of I/O Devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read-write	CD-ROM graphics controller disk

# Major Device Access Conventions

- Block I/O
- Character-stream I/O
- Memory-mapped file access
- Network sockets
- Clock and Time

# Outline

## 2 Application I/O Interface

- Block and Character Devices
- Network Devices
- Clocks and Timers
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

# Block and Character Devices

- 1 Block devices include disk drives
  - ▶ Commands include read, write, seek
  - ▶ Raw I/O or file-system access
  - ▶ Memory-mapped file access possible
- 2 Character devices include keyboards, mice, serial ports
  - ▶ Commands include get, put
  - ▶ **Libraries** layered on top allow line editing



# Outline

## 2 Application I/O Interface

- Block and Character Devices
- **Network Devices**
- Clocks and Timers
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

# Network Devices

- Varying enough from block and character to have own interface
- Unix and Windows NT/9x/2000 include **socket** interface
  - ▶ Separates network protocol from network operation
  - ▶ Server – socket, bind, listen, accept
  - ▶ Client – socket, connect
  - ▶ Includes select functionality
- Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)

# Outline

## 2 Application I/O Interface

- Block and Character Devices
- Network Devices
- **Clocks and Timers**
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

# Clocks and Timers

- Provide current time, elapsed time, timer
- **Hardware** clocks
  - 1 **Real Time Clock (RTC, 实时时钟)**
  - 2 **Time Stamp Counter (TSC, 时间戳计数器)**
  - 3 **Programmable Interval Timer (PIT, 可编程间隔定时器)**
    - ★ used for timings, periodic interrupts
- ioctl (on UNIX) covers odd aspects of I/O such as clocks and timers

# Clocks and Timers

## ① Real Time Clock (RTC, 实时时钟)

- ▶ Integrated with CMOS RAM, always tick.
- ▶ Seconds from **00:00:00 January 1, 1970 UTC**
- ▶ Can be used as an alarm clock
  - ★ IRQ8
  - ★ Interrupt frequency: 2HZ~8192HZ
- ▶ I/O address (port no): 0x70, 0x71
- ▶ Example:
  - ★ Motorola 146818: CMOS RAM + RTC
- ▶ Second ↔ year, month, date, week **HOW?**

# Clocks and Timers

## ② Time Stamp Counter (TSC, 时间戳计数器)

- ▶ 64bit TSC register in the processor
  - ★ Pentium and after
- ▶ Incremented at each clock signal on **CLK** input pin
  - ★ example: CPU frequency 400MHz  
adds 1 per 2.5 ns = adds  $400 \times 10^6$  per second
- ▶ Instruction: rdtsc
- ▶ How to know CPU frequency?

# Clocks and Timers

## ③ Programmable Interval Timer (PIT, 可编程间隔定时器)

- ▶ 8253, 8254
- ▶ Issues **time interrupt** in a programmable time interval
- ▶ Can also be used to calculate processor frequency during boot up.
- ▶ 8253
  - ★ 14,3178 MHz crystal  $\Rightarrow$  4,772,727 Hz system clock  $\Rightarrow$  1,193,180 Hz to 8253
  - ★ using 16 bit divisor  $\Rightarrow$  interrupt every 838 ns  $\sim$  54.925493 ms

# Outline

## 2 Application I/O Interface

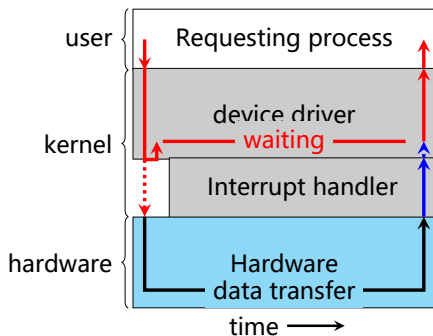
- Block and Character Devices
- Network Devices
- Clocks and Timers
- Blocking (阻塞) and Nonblocking (非阻塞) I/O



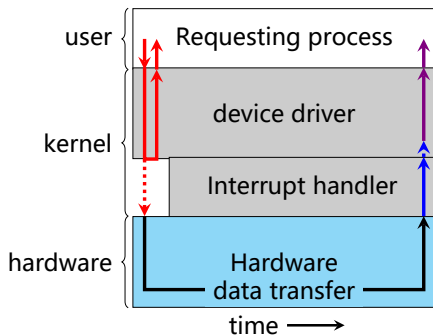
# Blocking (阻塞) and Nonblocking (非阻塞) I/O

- **Blocking** (阻塞) – process suspended until I/O completed
  - ▶ Easy to use and understand
  - ▶ Insufficient for some needs
- **Nonblocking** (非阻塞) – I/O call returns as much as available
  - ▶ User interface, data copy (buffered I/O)
  - ▶ Implemented via **multi-threading**
  - ▶ Returns quickly with count of bytes read or written
  - ▶ **Asynchronous** (异步) – process runs while I/O executes
    - ★ Difficult to use
    - ★ I/O subsystem signals process when I/O completed

# Two I/O Methods



(a)  
Synchronous



(b)  
Asynchronous

# Outline

## 3 Kernel I/O Subsystem

- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- Error Handling
- I/O Protection
- Kernel Data Structures

# Kernel I/O Subsystem Services

- Kernel I/O Subsystem Services

- ① I/O Scheduling
- ② Buffering
- ③ Caching
- ④ Spooling
- ⑤ Device reservation
- ⑥ Error handling

# Outline

## 3 Kernel I/O Subsystem

- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- Error Handling
- I/O Protection
- Kernel Data Structures

# I/O Scheduling

- **I/O scheduling:**

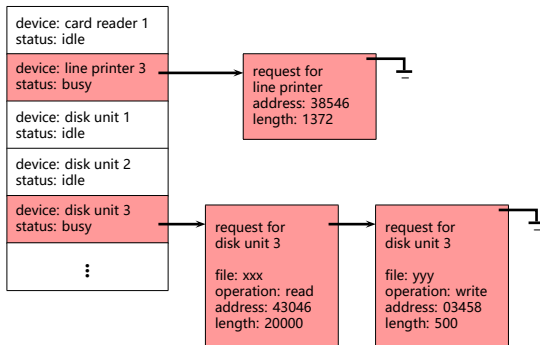
To **schedule** a set of I/O requests **means** to **determine a good order** in which to execute them

- ▶ **Origin order:** the order in which applications issue system calls:  
**May NOT the best order!**
- ▶ Scheduling can
  - ★ **Improve overall system performance**
  - ★ Share device access **fairly** among processes
  - ★ **Reduce the average waiting time** for I/O to complete
- ▶ Example: Disk read request from Apps.  
App1: 0; App2: 100; App3: 50;  
Now at 100;  
The OS may serve the applications in the order App2, App3, App1.

# I/O Scheduling

- OS maintaining a wait queue of request for each device

- **Device-status Table**



- I/O scheduling,  
Some OSes try fairness, some not

# I/O Scheduling

- Another way to improve performance is by using storage space in main memory or on disk
  - ▶ Buffering (缓冲机制)
  - ▶ Caching
  - ▶ Spooling



# Outline

## 3 Kernel I/O Subsystem

- I/O Scheduling
- **Buffering (缓冲机制)**
- Caching, Spooling & device reservation
- Error Handling
- I/O Protection
- Kernel Data Structures

# Buffering (缓冲机制)

- **Buffering (缓冲机制)**

- ▶ Buffer – A memory area that stores data while they are transferred between two devices or between a device and an application
- ▶ Store data in memory while transferring between devices

- Why buffering?

- ① To cope with device **speed** mismatch.

Example: Receive a file via modem and store the file to local hard disk.

- ★ Speed: The modem is about a thousand times slower than the hard disk.
- ★ Two buffers are used.

# Buffering (缓冲机制)

- **Buffering (缓冲机制)**

- ▶ Buffer – A memory area that stores data while they are transferred between two devices or between a device and an application
- ▶ Store data in memory while transferring between devices

- Why buffering?

- ② To cope with device transfer **size** mismatch.

Example: Send/receive a large message via network.

- ★ At sending side: the large message is fragmented into small network packets.
- ★ At receiving side: the network packets are placed in a reassembly buffer.

# Buffering (缓冲机制)

- **Buffering (缓冲机制)**

- ▶ Buffer – A memory area that stores data while they are transferred between two devices or between a device and an application
- ▶ Store data in memory while transferring between devices

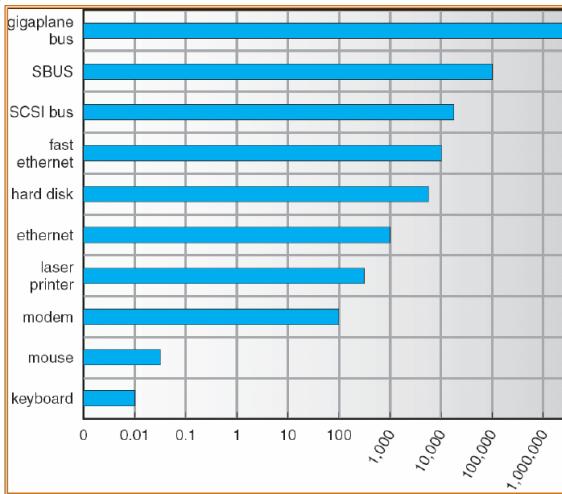
- Why buffering?

- ③ To maintain “copy semantics”

Example: When write() data to disk, it first copy the data from application' s buffer to a kernel buffer.

# Buffering (缓冲机制)

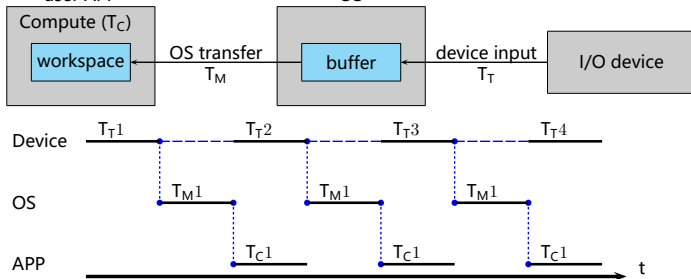
- Sun Enterprise 6000 Device-Transfer Rates



# Buffering (缓冲机制)

## ① Single buffer (单缓冲)

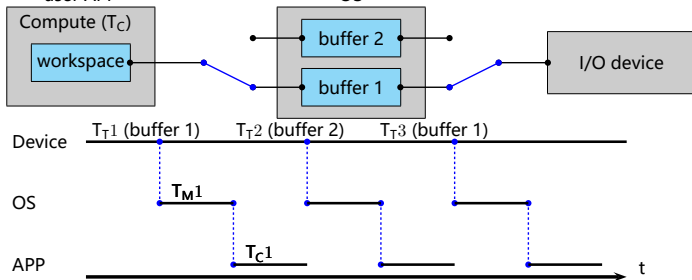
- ▶ APP.workspace  $\xleftrightarrow{(\text{OS}, T_M)}$  OS.buffer  $\xleftrightarrow{(\text{Device}, T_T)}$  Device
- ▶ Suppose the computing time of APP is  $T_C$ ,  
if current  $T_C$  can parallel with the next  $T_T$ ,  
we have  $T_{\text{average}} = \max(T_C, T_T) + T_M$



# Buffering (缓冲机制)

## ② Double buffer (双缓冲)

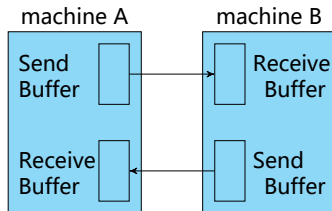
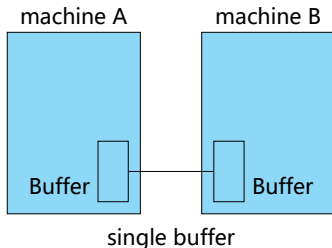
- $\approx \max(T_{C, T_T})$ ; 连续输入 ( $T_C < T_T$ ) 或者连续计算 ( $T_C > T_T$ )



# Buffering (缓冲机制)

## ② Double buffer (双缓冲)

- ▶ Another usage of single buffer and double buffers: in communication between two machines





# Buffering (缓冲机制)

## ③ Circular buffer (循环缓冲)

- ▶ Multiple (types of) buffers + multiple buffer pointers
  - ★ Empty buffers and  $Next_i$ ;  
Full buffers and  $Next_g$ ;  
the current buffer in consumption
- ▶ Similar to the PC problem.

## ④ Buffer pool (缓冲池)

- ▶ 前三种，缓冲区是专用的
- ▶ 为提高缓冲区利用率：设置公共的缓冲池

# Outline

## 3 Kernel I/O Subsystem

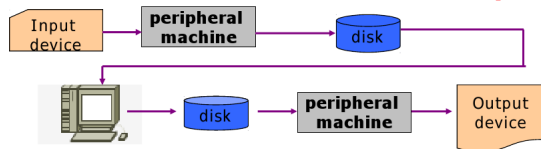
- I/O Scheduling
- Buffering (缓冲机制)
- **Caching, Spooling & device reservation**
- Error Handling
- I/O Protection
- Kernel Data Structures

# Caching, Spooling & device reservation

- ① **Caching** - fast memory holding copy of data
  - ▶ Always just a copy
  - ▶ Key to performance
- ② **Spooling** - hold output for a device
  - ▶ **Dedicated device** can serve only one request at a time
  - ▶ Spooling is a way of dealing with I/O devices in a multiprogramming system
  - ▶ Example: Printing
- ③ **Device reservation** - provides exclusive access to a device
  - ▶ System calls for allocation and deallocation
  - ▶ Watch out for deadlock

# Spooling

- **Out-line I/O** (脱机I/O), 使用**外围机 (peripheral machine)**



- **SPOOL:**  
Simultaneous Peripheral Operation On-Line  
(外部设备联机并行操作, **假脱机**)
  - ▶ Dedicated device → sharable device
  - ▶ Using **processes** of multiprogramming system

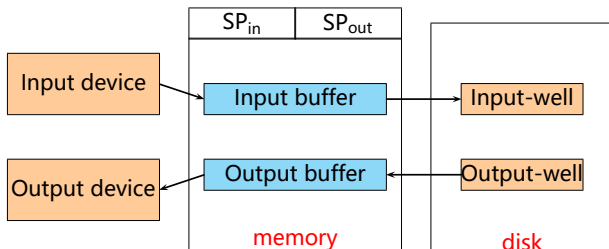
# Spooling

## ● SPOOL:

Simultaneous Peripheral Operation On-Line  
(外部设备联机并行操作, **假脱机**)

### ► Structure

- ★ Input-well (输入井), output-well (输出井)
- ★ Input-buffer, output-buffer
- ★ Input-process  $SP_{in}$ , output-process  $SP_{out}$
- ★ **Requested-queue**



# Outline

## 3 Kernel I/O Subsystem

- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- **Error Handling**
- I/O Protection
- Kernel Data Structures

# Error Handling

- OS can **recover** from disk read, device unavailable, transient write failures
  - ▶ Example: read() again, resend(), ..., according to some specified rules
- Most return an **error number** or code when I/O request fails
- **System error logs** hold problem reports

# Outline

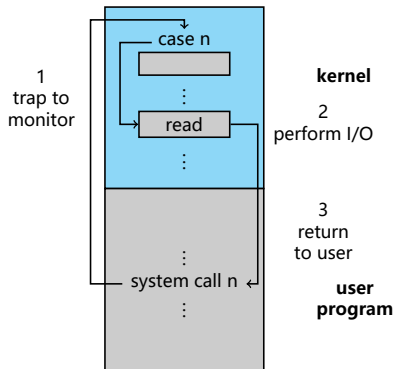
## 3 Kernel I/O Subsystem

- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- Error Handling
- **I/O Protection**
- Kernel Data Structures



# I/O Protection I

- User process may accidentally or purposefully attempt to disrupt normal operation via illegal I/O instructions
- **To prevent users from performing illegal I/O**
  - ▶ All I/O instructions defined to be **privileged**
  - ▶ I/O must be performed via **system calls**
    - ★ Memory-mapped and I/O port memory locations must be protected too



Use of a System Call to Perform I/O

# Outline

## 3 Kernel I/O Subsystem

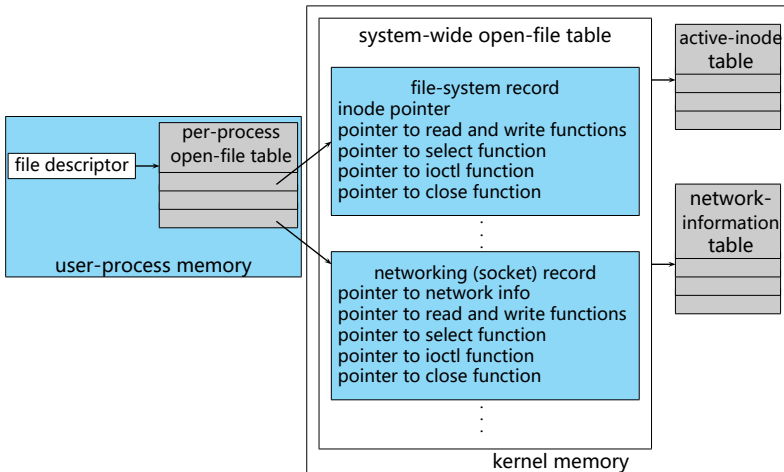
- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- Error Handling
- I/O Protection
- Kernel Data Structures

# Kernel Data Structures

- Kernel **keeps state** info for I/O components, including
  - ▶ open file tables,
  - ▶ network connections,
  - ▶ character device state
- Many, many complex data structures to **track buffers, memory allocation, “dirty” blocks**
- Some use object-oriented methods and message passing to implement I/O

# Kernel Data Structures

- Example: UNIX I/O Kernel Structure



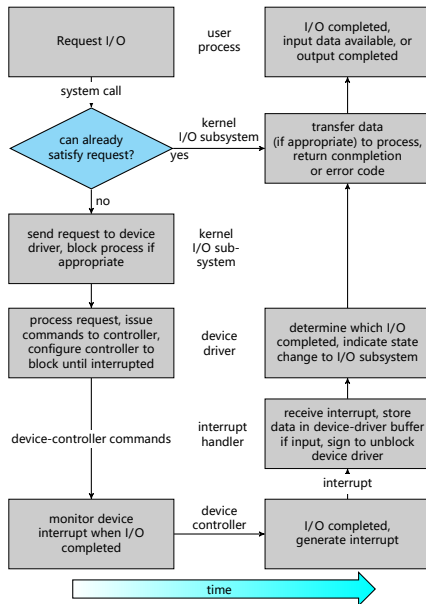
# Outline

## 4 Transforming I/O Requests to Hardware Operations

# I/O Requests to Hardware Operations

- Consider reading a file from disk for a process:
  - ① Determine device holding file
  - ② Translate name to device representation
  - ③ Physically read data from disk into buffer
  - ④ Make data available to requesting process
  - ⑤ Return control to process

# The Typical Life Cycle of An I/O Request



# Outline

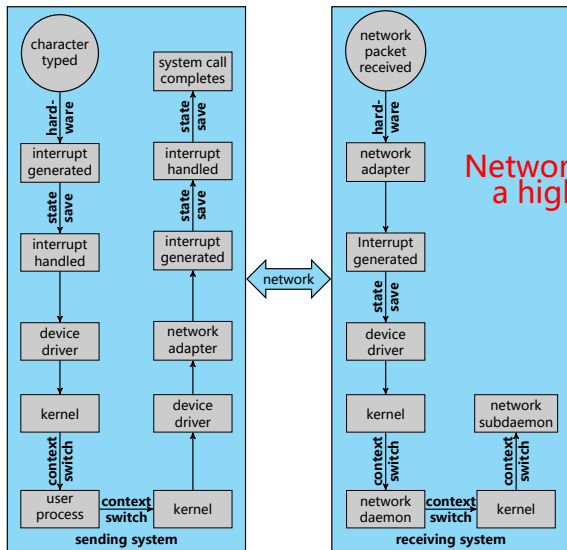
## 5 Performance



# Performance

- **I/O is a major factor in system performance:**
  - ▶ Demands CPU to execute device driver, kernel I/O code
  - ▶ Context switches due to interrupts
  - ▶ Data copying
  - ▶ Network traffic especially stressful

# Intercomputer Communications



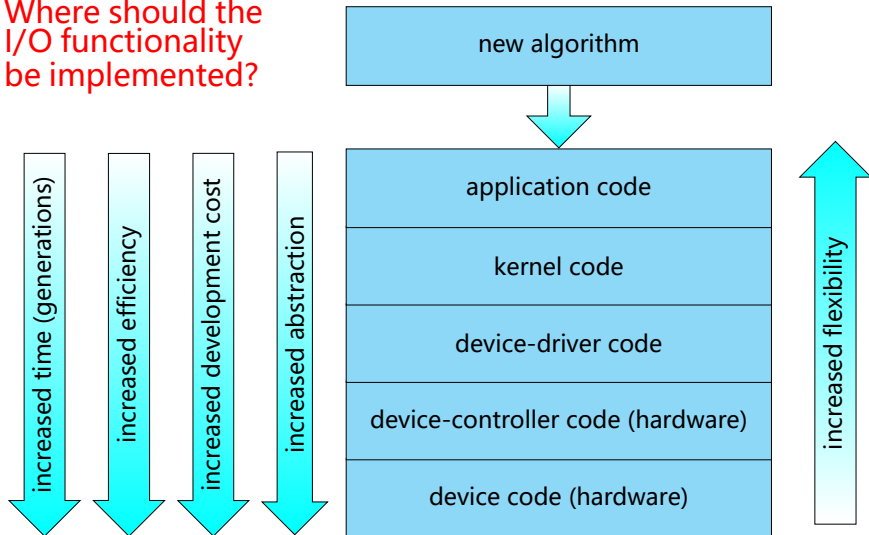
Network traffic can also cause a high context-switch rate

# Improving Performance

- 1 Reduce number of context switches
- 2 Reduce data copying
- 3 Reduce interrupts by using large transfers, smart controllers, polling
- 4 Use DMA
- 5 Move processing primitives into hardware
- 6 Balance CPU, memory, bus, and I/O performance for highest throughput

# Device-Functionality Progression

Where should the I/O functionality be implemented?



# Outline

## 6 小结

# 小结

## 1 I/O Hardware and I/O control methods

- Polling (轮询方式)
- Interrupts (中断方式)
- Direct Memory Access (DMA方式)
- I/O hardware summary

## 2 Application I/O Interface

- Block and Character Devices
- Network Devices
- Clocks and Timers
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

## 3 Kernel I/O Subsystem

- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- Error Handling
- I/O Protection
- Kernel Data Structures

## 4 Transforming I/O Requests to Hardware Operations